

Intants — Production Expansion Plan

Status: Planning document (v1, 2026-06-17) **Owner:** Intants engineering **Scope:** Evolving the current 1:1 voice-avatar interviewer into a production platform with candidate video, gaze-based proctoring, bulk hiring, multiple intake scenarios, and production-grade security. This document describes **what we will build and why**. It does not contain final code. Each phase links back to the real services it touches so the work is traceable. Read alongside HLD.md, LLD.md, and Final_stack.md.

1. Where we are today (baseline)

The platform is a **1:1 voice-first AI interview system** and it works end-to-end:

Capability	Where it lives	State
Real-time avatar interview (Tavus over LiveKit)	services/interview_core (app/worker/interview_worker.py)	■ Working
Voice pipeline (Sarvam STT/TTS, Silero VAD)	interview_core/app/speech/	■ Working
Interviewer brain / question policy	interview_core/app/worker/interview_worker.py (_interviewer_instructions)	■ Working
Resume-grounded questions	resume text injected into the interviewer prompt	■ Working
Auth + users + consent (pluggable)	services/data_gateway	■ Working
Scoring + PDF scorecard	services/feedback_billing	■ Working (needs Gemini key)
Admin dashboard + analytics	services/admin_ops + web/src/pages/admin/	■ Working

Two things the current product does NOT have, and this plan adds:

- **No candidate video / camera.** The candidate is audio-only; only the avatar has a video track. There is no webcam capture, no candidate-side video, no true two-way "face-to-face" feel.
 - **No proctoring / integrity signals.** There is no detection of looking away, absence, multiple faces, or tab-switching. Nothing prevents or flags malpractice.
- Everything below is **additive** and env-swappable, consistent with the two-tier strategy in Final_stack.md.

2. Goals of this expansion

- **Candidate video + true 1:1 interaction** — the candidate's webcam is live in the session (shown to themselves, captured for proctoring), alongside the avatar. The session feels face-to-face.
- **Gaze / eye-movement proctoring** — detect when the candidate looks away from the screen, leaves the frame, has multiple faces present, or switches tabs. Produce an **integrity score** for human review.
- **Bulk hiring at scale** — one or many companies running thousands of interviews in a hiring drive, with invites, batches, and a recruiter console.
- **Multiple intake scenarios** — off-campus single applicant, resume-driven screening, and fully admin-managed campaigns — all on one engine.

- **Production security** — the controls required before this touches real candidate PII and biometric data at scale (DPDP Act 2023 + RFP NFRs).

3. A note on terminology — "eye tracking" vs what is actually feasible

This matters for credibility, especially with government buyers.

- **True retina / iris tracking** (knowing the exact pixel a person is looking at) requires **infrared hardware** (e.g. Tobii). It is **not possible with a normal webcam**. We will not claim it.

- **What is feasible and is what real proctoring products ship:** webcam-based **gaze direction + head pose estimation** (approximate "is the candidate looking at the screen or away"), plus face-presence, multi-face, and browser-event signals. This is ~80–90% reliable and runs free, in the browser. We will frame the feature honestly as **"AI-assisted integrity flagging for human review"** — never "automatic cheating detection." False positives in an automated-reject system create legal and bias risk.

4. Phase A — Candidate video + true 1:1 interaction

Objective: the candidate's webcam is live in the interview session.

What changes

- **Frontend** (web/src/features/interview/): request **camera** permission (today we only request mic), render the candidate's own video in a corner / side panel next to the avatar, and **publish the camera track to the LiveKit room** (LiveKit already carries our audio + avatar video, so the candidate video track is the same transport — no new infra).
- **Consent** (data_gateway consent ledger): a ****new DPDP consent type** **video_capture**** must be granted before the camera turns on. This is our first candidate-side video, so it needs its own explicit consent record.
- **Storage decision (cost + DPDP):** by default ****do NOT** upload raw candidate video** to our servers. Keep it client-side and ephemeral. Recording the full video is a separate, opt-in, consented feature (large storage + biometric data under DPDP §3(k)). For most use cases, only the ***proctoring events*** (Phase B) need to leave the browser, not the video itself.

Effort

~2–3 days (frontend + consent type + a migration for the new consent type).

Why this is first

Phase B (gaze) needs a camera stream to analyze. Phase A delivers the camera and the consent gate that Phase B builds on.

5. Phase B — Gaze tracking & proctoring pipeline

Objective: detect malpractice signals and produce a reviewable integrity score.

5.1 Detection — runs in the browser, not on our servers

Detection runs **client-side in a Web Worker** so it costs us no server compute and the raw video never leaves the candidate's machine (cheapest + safest for DPDP). The browser emits small **events**, not video frames.

Signal	Method	Reliability
Gaze direction / "looking away"	MediaPipe FaceLandmarker (head pose + eye landmarks)	~80–85%
Face present / absent	Same model (face count = 0)	~95%
Multiple faces in frame	Same model (face count > 1)	~98%
Tab / window switch, focus loss	Browser <code>visibilitychange</code> + <code>blur</code>	~99.9%
Copy / paste, fullscreen exit	Browser clipboard / fullscreen API	~99.9%
Second voice / background speech	Reuse existing Silero VAD in <code>interview_core</code>	~90%

Recommended model stack (all free, browser-based):

- **MediaPipe Tasks (FaceLandmarker)** — primary. 478 landmarks, head pose, eye region. ~30ms/frame, no calibration. @mediapipe/tasks-vision.

- **MediaPipe Iris** — optional add-on if we later want finer eye precision.

- **WebGazer.js** — only if we want true on-screen gaze coordinates (needs a ~30s 9-point calibration; more UX friction). Not needed for v1.

- A **Python / server-side** path (e.g. MediaPipe Python, OpenCV) is the **Tier-2

option if we ever move detection server-side for tamper-resistance — but that adds GPU/CPU cost per concurrent session and re-introduces the DPDP video- transit problem. Client-side first.**

*Note on the CLAUDE.md "Python models" idea: gaze CAN be done with Python (MediaPipe Python), but doing it **server-side** means streaming each candidate's video to us — expensive at 10k concurrent and a DPDP liability. The browser (JS/WASM MediaPipe) is the right default; keep Python server-side detection as a Tier-2 tamper-resistance upgrade, gated on cost review.*

5.2 Event flow & scoring

- Web Worker runs inference every ~500ms.

- It emits lightweight events: `gaze_away_start/end`, `face_absent_start/end`, `multiple_faces`, `tab_blur/focus`, `second_voice`.

- Events are batched and POSTed to a new endpoint in `interview_core`.

- Events are stored against the session (new table — see 5.3).

- At session end, a **scorer** computes an **integrity score (0–100)** from time-in-flagged-states + event counts, plus a timeline.

- The score + timeline surface in the **admin dashboard** drill-in

(`web/src/pages/admin/AdminInterviewDetail.tsx`) for **human review**.

5.3 New data model

```
integrity_events
id uuid pk
session_id uuid fk -> sessions(id) on delete cascade
event_type text -- gaze_away | face_absent | multiple_faces | tab_blur | second_voice | ...
started_at timestamptz
ended_at timestamptz null -- for ranged events
metadata jsonb -- e.g. {"confidence": 0.7, "yaw_deg": 35}

sessions (add columns)
integrity_score smallint null -- 0-100, computed at session end
```

```
proctoring_summary jsonb null -- counts + total flagged seconds
```

One Alembic migration in `data_gateway`. `interview_core` already maps a partial sessions model; add the two columns there too (read path).

5.4 Effort

~5–7 days: Web Worker + MediaPipe integration (2), event API + schema (1), integrity scoring (1), admin timeline UI (1–2), consent + tests (1).

5.5 Honesty / fairness guardrails (do not skip)

- Label everything "flag for review," never "cheating confirmed."
- Calibrate thresholds; webcam gaze is noisy (glasses, lighting, dark skin tones, disability). Document known limitations.
- Make the candidate aware proctoring is active (consent + visible indicator).

6. Phase C — Bulk hiring engine

Objective: support a company (or many) interviewing thousands of candidates in a hiring drive.

*Architectural reality check (decided 2026-06-16): the chosen direction is **avatar-everywhere + fully-synchronous live** interviews. That is the most demanding and most expensive mode. At thousands of concurrent live avatar sessions we are bound by **Tavus/Simli vendor concurrency limits** (needs an enterprise quota, not the demo tier) and very likely exceed the **12/session cap** in `Final_stack.md`. This must be validated by cfo-cost-watcher + a real load test (the unused Locust scaffold, backlog B-029) **before** any government bid. Treat full-sync-avatar as the premium tier; consider voice-only or async for true mass screening if economics demand it.*

6.1 New domain model — campaigns & batches

The current model is 1:1 (one user, one ad-hoc session). Bulk hiring needs a campaign concept owned by a company/recruiter:

```
companies
id, name, ... -- the hiring org (or college)

campaigns
id, company_id, job_id, name
window_start, window_end -- when candidates may interview
settings jsonb -- avatar on/off, proctoring level, language, question count
status -- draft | open | closed

campaign_invites
id, campaign_id, candidate_email
token -- unique magic-link token
status -- invited | started | completed | expired
session_id uuid null fk -> sessions

sessions (add column)
campaign_id uuid null fk -> campaigns(id) -- null = self-serve / off-campus
```

6.2 New infrastructure — a real job queue

We have **no job queue today** (everything is synchronous FastAPI). Bulk needs one:

- **Recommended:** Arq or Celery on the **Upstash Redis** we already run.

- Queued jobs: send invites + reminder emails (Resend is in-stack), end-of-session

scoring, PDF generation, post-campaign report aggregation.

6.3 Worker pool & autoscaling

- The `interview_worker` (LiveKit agent) becomes a **horizontally-scaled pool**.

Each live interview = one agent process. The pool size is capped by the avatar vendor concurrency quota.

- On Tier-2 (AWS EKS Mumbai), this is an HPA-scaled deployment; on the demo tier (Railway), it's manual replica scaling.

6.4 Recruiter console (new frontend area)

- CSV / ATS upload of candidate lists → generates `campaign_invites`.
- Live campaign dashboard: X invited / Y started / Z completed / avg score.
- Bulk export of scorecards.
- Lives alongside `web/src/pages/admin/` (reuses the `admin_ops` analytics patterns).

6.5 Effort

~3–4 weeks (campaign model + migrations, job queue, invite flow + emails, recruiter console, worker autoscaling, load test).

7. Intake scenarios (one engine, several front doors)

All scenarios drive the **same** interview engine; they differ only in how a session is created and what context is attached.

#	Scenario	How a session is created	Notes
1	Off-campus individual	Candidate self-registers, picks a role, starts	Today's flow. <code>campaign_id = null</code> .
2	Resume-driven	Candidate uploads resume → questions grounded in it	Already wired (<code>resume_text</code> → interviewer prompt). Future: <code>resume</code> ↔ JD fit score as a pre-gate.
3	Bulk / campaign	Recruiter invites via campaign; candidate clicks magic link	Phase C. <code>campaign_id</code> set; settings inherited from campaign.
4	Admin-managed	Admin creates jobs/JDs, configures campaigns, reviews results	<code>admin_ops</code> + recruiter console; admin role (see <code>grant_admin.py</code>).

The key design principle: **scenario = a thin layer over session creation**, not a fork of the interview engine. This keeps one codebase for all markets.

8. Production security (required before real PII at scale)

Maps to DPDP Act 2023 + the RFP NFRs in `HLD.md`. Several items already exist; this lists what must be true for production.

Already in place

- Pluggable auth, JWT with issuer/audience validation, role-gated admin.
- DPDP consent ledger, right-to-erasure pipeline, audit log, soft-delete + retention.

- Secrets via env / Pydantic settings (no hardcoded secrets).

Must add / harden for production

- **Biometric data governance (NEW, critical):** camera + gaze data are

biometric under DPDP §3(k). Needs: explicit video_capture / proctoring consent types, a defined retention window, encryption at rest if any video is ever stored, and inclusion in the erasure pipeline.

- **India data residency:** Tier-2 production = AWS Mumbai only (already the

plan). Tavus is US-hosted and demo-only — must be replaced by the custom RPM avatar (Tier-2) before a government deployment.

- **Rate limiting & abuse controls:** invite-token brute-force protection,

per-IP and per-account limits (login limiter exists; extend to invites/sessions).

- **Transport & secrets:** TLS everywhere, secrets in a manager (AWS Secrets

Manager on Tier-2), rotate the demo keys before production.

- **Proctoring fairness & transparency:** documented accuracy limits, human-in-

the-loop review, candidate notification — to defend against bias/legal claims.

- **Tenant isolation (multi-company):** campaign/company data must be scoped so

one company can never read another's candidates or results.

- **Sign-off gates:** security-auditor review before any production deploy;

cfo-cost-watcher review of per-session cost before any L1 bid.

9. Recommended sequencing

Order	Phase	Why	Rough effort
1	A — Candidate video	Unlocks B; small, high-impact	2–3 days
2	B — Gaze proctoring	Biggest differentiator; self-contained	5–7 days
3	C — Bulk engine	The scale story; largest build	3–4 weeks
4	Security hardening	Continuous; gate before production	ongoing

Phases A + B can ship on the current architecture without the bulk engine, so the proctoring story is demoable long before the scale work lands.

10. Open decisions to confirm before building

- **Record candidate video, or events-only?** (Events-only = cheaper + DPDP-safe;

recording = stronger evidence but storage + biometric burden.)

- **Bulk mode economics:** confirm avatar-everywhere full-sync is financially

viable at scale, or define a voice-only/async fallback for mass screening.

- **Proctoring strictness:** advisory flags only (recommended) vs. hard auto-

actions (not recommended for v1).

- **Multi-company now or single-tenant first?** (Affects whether we build the

companies table + tenant isolation in Phase C or defer it.)

End of plan. Update this document as decisions are made; it is the source of truth for the expansion scope.